

Weifan Lin

Computer Science

Csc342 Section G

04/23/2015

Take-Home Exam

Title:

Matrix Multiplication

Objective:

The goal of this take-home exam was to implement matrix multiplication in different ways and to compare their performance. Firstly to create a simple c function to compute matrix-matrix multiplication, then to use DPPS instruction to improve performance. In addition, I also had some functions to compute the time of different performances for time analysis at the end.

C++ function to compute matrix multiplication

```
57 // function written in C++ language
58 float* MatrixVector_C(float* matrix, float* vector, int size)
59 {
60     float *vectorResult=(float*) malloc (size*sizeof(float));
61     float* matrix_ptr=matrix;
62
63     for (int i=0; i<size; i++)
64     {
65         float* vector_ptr=vector;
66
67         float Result=0.0;
68         for (int j=0; j<size; j++)
69         {
70             Result+=(*matrix_ptr)*(*vector_ptr);
71             vector_ptr++;
72             matrix_ptr++;
73         }
74         vectorResult[i]=Result;
75     }
76     return (vectorResult);
77 }
```

Use DPSS instruction to Improve performance

```
// function written in assembler (SSE)
float* MatrixVector_SSE(float* refmatrixA,float* refvectorB, int size)
{
    int _length=size;
    int length1=size;
    int _nCol=size;
    int _nRow=size;

    float *vect=refvectorB;
    float *vect1=(float*) malloc (size*sizeof(float));
    float *matr=refmatrixA;

    int stride1=_nCol<<2;
    int stride2=stride1<<1;
    int stride3=stride2+stride1;
    int stride4=_nCol<<4;

    int nbre=_nRow>>2;
```

```

int nbre1=_nCol>>2;
int nbre2=nbre1;

int resteRow=_nRow%4;
int resteCol=_nCol%4;
int nRow1=_nRow;
int nRow2=_nRow;
int nCol1=_nCol;

__declspec(align(16)) float valeur1[4];
__declspec(align(16)) float valeur2[4];
__declspec(align(16)) float valeur3[4];
__declspec(align(16)) float valeur4[4];

float* derniersLigne=matr+(stride1>>2)-resteCol;
float* derniersVect=vect-resteCol+length1;
float zero=0;
int stride=stride1-(resteCol<<2);
nRow1-=resteRow;

// Is the matrix dimension bigger than 4x4?
__asm
{
    push    ebx           ; save EBX (warning :frame pointer register)
                        ; we save the data in EBX to avoid bugs due to warnings C4731

    // Is the matrix dimension bigger than 3?
    cmp     _nCol, 3
    jle     inferieur
    cmp     _nRow, 3
    jle     inferieur

    /* Init of the result vector (all values to zero) */
    xorps   xmm0,xmm0
    mov     edi,vect1
    mov     eax,nbre
    mov     ecx, resteRow

init:  movups [edi],xmm0
    add     edi,16
    dec     eax
    jnz     init

    cmp     ecx, 0
    je      debut
    // Init the end of the result vector if its size is not a multiple of 4

```

```

init_fin:
    movss [edi], xmm0
    add    edi, 4
    dec    ecx
    jnz    init_fin

```

```

debut:
    // Init vectors' pointers
    mov     esi, vect
    mov     edi, vect1

    // Init matrix pointer and strides
    mov     ebx, matr
    mov     eax, stride1
    mov     ecx, stride2
    mov     edx, stride3

```

```

// the matrix is cutted in little matrix 4x4
// and the vector in little vector 4x1

```

```

ligne: push    ebx    ; save ebx in the stack
        push    nbre   ; save nbre in the stack
        push    edi    ; save edi in the stack

```

```

    movups xmm0, [esi]

```

```

    movaps xmm1, xmm0
    movaps xmm2, xmm0
    movaps xmm3, xmm0
    shufps xmm0, xmm0, 0x00
    shufps xmm1, xmm1, 0x55
    shufps xmm2, xmm2, 0xAA
    shufps xmm3, xmm3, 0xFF

```

```

    movaps valeur1, xmm0
    movaps valeur2, xmm1
    movaps valeur3, xmm2
    movaps valeur4, xmm3

```

```

// the vector 4x1 is loaded in the XMM0,XMM1,XMM2,XMM3

```

```

colonne:

```

```

    // Now we transpose the matrix in order to perform operations

```

```

    movlps xmm4, [ebx]

```

```
movlps xmm6, [ebx+0x08]
movlps xmm3, [ebx+ecx] // ->2*stride
movlps xmm2, [ebx+ecx+8] // ->2*stride+8
```

```
movhps xmm4, [ebx+eax] // ->stride
movhps xmm6, [ebx+eax+8] // ->stride+8
movhps xmm3, [ebx+edx] // ->3*stride
movhps xmm2, [ebx+edx+8] // ->3*stride+8
```

```
movaps xmm5, xmm4
movaps xmm7, xmm6
shufps xmm4, xmm3, 0x88
shufps xmm6, xmm2, 0x88
```

```
shufps xmm5, xmm3, 0xDD
shufps xmm7, xmm2, 0xDD
```

```
// we have the matrix transpose in registers XMM4 to XMM7
```

```
// load vector 4x1 values in XMM0 to XMM3
```

```
movaps xmm0, valeur1
movaps xmm1, valeur2
movaps xmm2, valeur3
movaps xmm3, valeur4
```

```
// perform multiplication between matrix values and vector values
```

```
mulps xmm0, xmm4
mulps xmm1, xmm5
mulps xmm2, xmm6
mulps xmm3, xmm7
```

```
// perform addition
```

```
addps xmm1, xmm0
addps xmm2, xmm3
addps xmm1, xmm2
```

```
movups xmm2, [edi]
addps xmm1, xmm2
```

```
/* the result (XMM1) is saved in the result vector */
```

```
movups [edi], xmm1
```

```
// Next little matrix 4x4 and next vector 4x1
```

```
add ebx, stride4
```

```
add    edi, 16
dec     nbre
jnz     colonne
```

```
pop     edi
pop     nbre
pop     ebx
add     ebx, 16
add     esi, 16
```

```
dec     nbre1
jnz     ligne
```

```
// This section of the function performs the last values of a row
// when the matrix dimension is not a multiple of 4
```

```
cmp     resteCol, 0
jz      fin
mov     eax, nRow2
mov     esi, derniersLigne
mov     edi, vect1
```

ligne_suivante:

```
mov     ecx, resteCol
mov     ebx, derniersVect
```

```
movss   xmm2, zero
```

fin_ligne:

```
movss   xmm0, [esi]
movss   xmm1, [ebx]
mulss   xmm0, xmm1
add     esi, 4
add     ebx, 4
addss   xmm2, xmm0
dec     ecx
jnz     fin_ligne
movss   xmm3, [edi]
addss   xmm3, xmm2
movss   [edi], xmm3
add     edi, 4
add     esi, stride
dec     eax
jnz     ligne_suivante
```

```
fin:    cmp     resteRow, 0
je      end
```

```
}
```

```
    // Pointers used by the function to perform the last coefficients of the result vector  
    // used when the matrix dimension is not a multiple of 4  
float* dernieres_lignes=matr+(nRow2-resteRow)*nCol1;  
float* derniers_coeff=vect1+nRow2-resteRow;  
int resteCol4=resteCol<<2;
```

```
__asm  
{  
    mov     esi, dernieres_lignes  
    mov     ebx, vect  
    mov     edi, derniers_coeff  
    mov     eax, resteRow  
    mov     ecx, nbre2
```

```
debut_ligne:  
    movss   xmm2, zero
```

```
prochains_coeff:
```

```
    movups  xmm0, [esi]  
    movups  xmm1, [ebx]  
    mulps   xmm0, xmm1  
    movaps  xmm3, xmm0  
    movaps  xmm4, xmm0  
    movaps  xmm5, xmm0  
    shufps  xmm3, xmm0, 0x55  
    shufps  xmm4, xmm0, 0xAA  
    shufps  xmm5, xmm0, 0xFF  
    addps   xmm0, xmm3  
    addps   xmm0, xmm4  
    addps   xmm0, xmm5  
    addps   xmm2, xmm0  
    add     esi, 16  
    add     ebx, 16  
    dec     ecx  
    jnz     prochains_coeff  
    movss   xmm7, [edi]  
    addss   xmm2, xmm7  
    movss   [edi], xmm2  
    add     edi, 4  
    add     esi, resteCol4  
    mov     ebx, vect  
    mov     ecx, nbre2  
    dec     eax
```

```

        jnz    debut_ligne

    }

// When the matrix dimension is lower than 4

    _asm
    {
        jmp    end

inferieur:
    mov     esi, matr
    mov     edi, vect1
    mov     ebx, vect
    mov     eax, nRow2
    mov     edx, nCol1

mise_zero:
    movss   xmm2, zero

lignes:
    movss   xmm0, [esi]
    movss   xmm1, [ebx]
    mulss   xmm0, xmm1
    add     ebx, 4
    add     esi, 4
    addss   xmm2, xmm0
    dec     edx
    jnz     lignes
    movss   [edi], xmm2
    add     edi, 4
    mov     ebx, vect
    mov     edx, nCol1
    dec     eax
    jnz     mise_zero
end:
    pop     ebx
    }
    return (vect1);
}

```


Functions For Time Analysis:

```
// Read the computer's timer RDTSC
__int64 GetTime()
{
    __int64 clock;

    __asm
    {
        rdtsc          // Resad the RDTSC Timer
        mov     dword ptr[clock], eax    // Store the value in EAX and EDX
        mov     dword ptr[clock+4], edx

    }
    return clock;
}

// Perform and display the time saved or lost between the C++ and assembly
void TimeImprove(__int64 timeC, __int64 timeSSE)
{
    float gain=0;

    if(timeC>timeSSE)
    {
        gain=(1-((float)timeSSE/((float)timeC))*100;
        printf("\nTime saved: %f %% \n",gain);
    }
    if(timeC==timeSSE)
    {
        printf("\nTemps saved: 0 %%\n");
    }
    if(timeC<timeSSE)
    {
        gain=((float)timeSSE/((float)timeC)*100-100;
        printf("\nTemps lost: %f %% \n ", gain);
    }
}
```

Comparison

```
// The main contain the call to the functions, declarations, and time measurement
int main(int argc, char* argv[])
{
    //Enter the size of the matrix
    int size = 1024;
    int i;

    // Allocate memory
    float* matrix=(float*) malloc (size*size*sizeof(float));
    float* vector=(float*) malloc (size*sizeof(float));
    float* result=(float*) malloc (size*sizeof(float));

    float* matrix1=(float*) malloc (size*size*sizeof(float));
    for(i=0; i<size*size; i++)
        matrix1[i]=(float)i;

    // Writing values in the matrix and vector
    MatrixVectorWriting(matrix, vector, size);

    // Benchmark the two functions
    __int64 t1=GetTime();
    for(i=0; i<100; i++)
        result=MatrixVector_C(matrix,vector, size);
    __int64 t2=GetTime();
    __int64 time_C=t2-t1;
    printf("Time spend in C++ function: %d clock cycles.\n",time_C);

    __int64 t3=GetTime();
    for(i=0; i<100; i++)
        result=MatrixVector_SSE(matrix1, vector, size);
    __int64 t4=GetTime();
    __int64 time_SSE=t4-t3;
    printf("Time spend in Asm SSE function: %d clock cycles.\n",time_SSE);

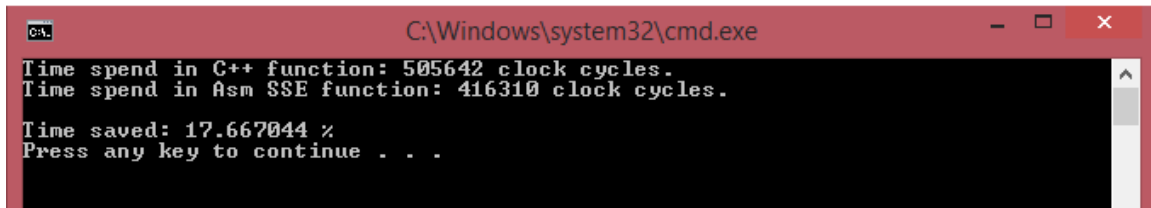
    // Display the time improvement in percent
    TimeImprove(time_C,time_SSE);
    system("pause");
    return 0;
}
```

The variable size in the main function is the size of matrix, we can change it manually to compare different performances of matrix multiplication in different sizes.

Time Analysis

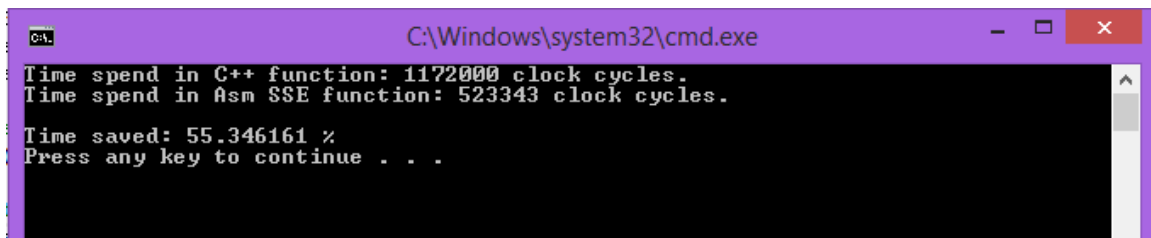
I set the size of matrix from 8x8 to 512x512, below are result of each one of them

size=8x8:



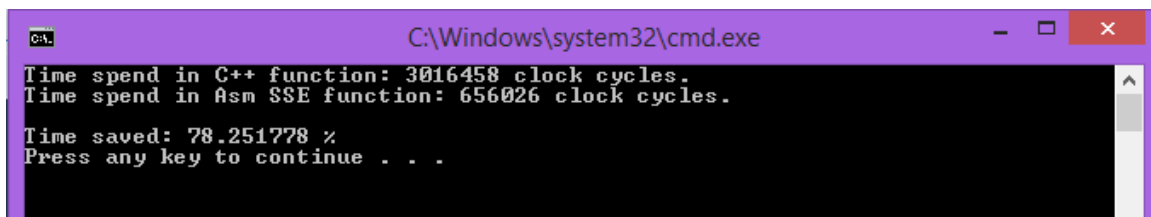
```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 505642 clock cycles.
Time spend in Asm SSE function: 416310 clock cycles.
Time saved: 17.667044 %
Press any key to continue . . .
```

size=16x16:



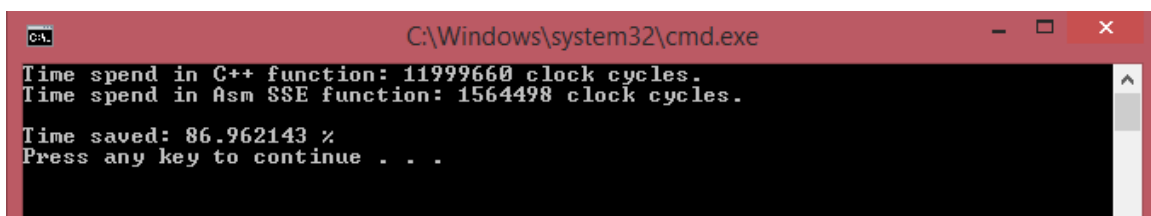
```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 1172000 clock cycles.
Time spend in Asm SSE function: 523343 clock cycles.
Time saved: 55.346161 %
Press any key to continue . . .
```

size=32x32:



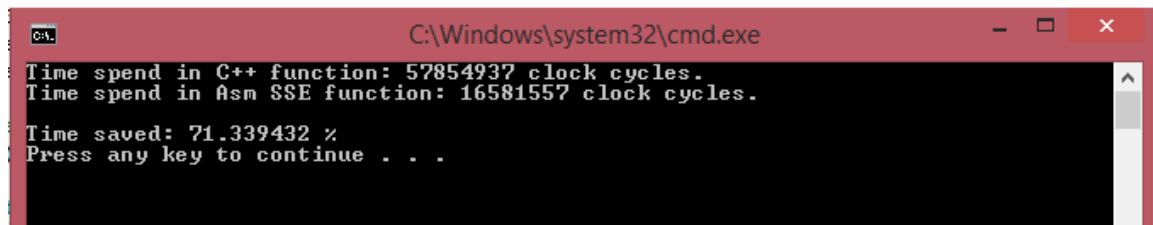
```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 3016458 clock cycles.
Time spend in Asm SSE function: 656026 clock cycles.
Time saved: 78.251778 %
Press any key to continue . . .
```

size=64x64:



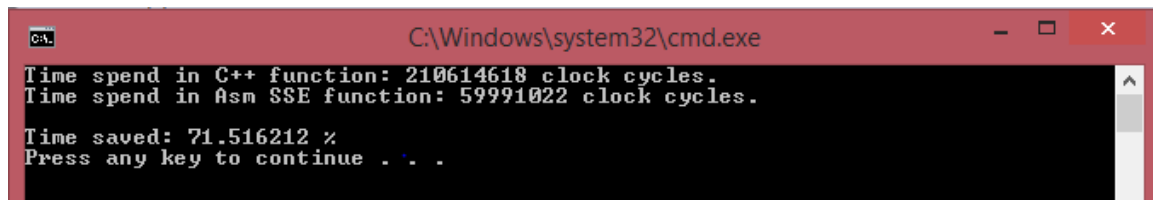
```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 11999660 clock cycles.
Time spend in Asm SSE function: 1564498 clock cycles.
Time saved: 86.962143 %
Press any key to continue . . .
```

size=128x128:



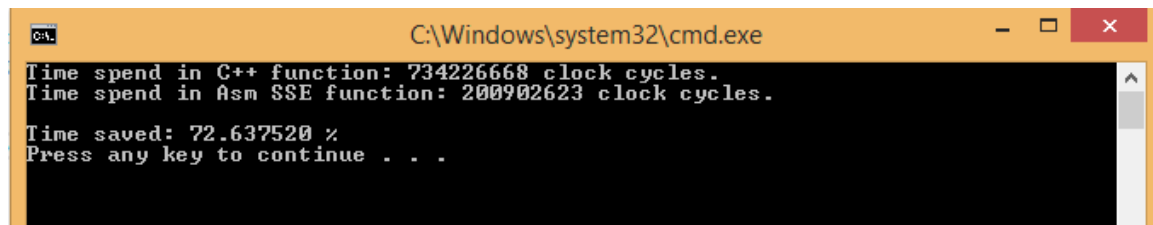
```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 57854937 clock cycles.
Time spend in Asm SSE function: 16581557 clock cycles.
Time saved: 71.339432 %
Press any key to continue . . .
```

size=256x256:



```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 210614618 clock cycles.
Time spend in Asm SSE function: 59991022 clock cycles.
Time saved: 71.516212 %
Press any key to continue . . .
```

size=512x512:



```
C:\Windows\system32\cmd.exe
Time spend in C++ function: 734226668 clock cycles.
Time spend in Asm SSE function: 200902623 clock cycles.
Time saved: 72.637520 %
Press any key to continue . . .
```

Conclusion

As we can see the results from the time analysis, I have shown that for a large number of repeated operations, the use of vector processing is much faster than single instruction processing. Overall, this take-home exam helped me implement matrix multiplication in

single instruction processing and in vector processing. Through the experiment I saw the performance of vector processing is faster as the size of matrix increases.